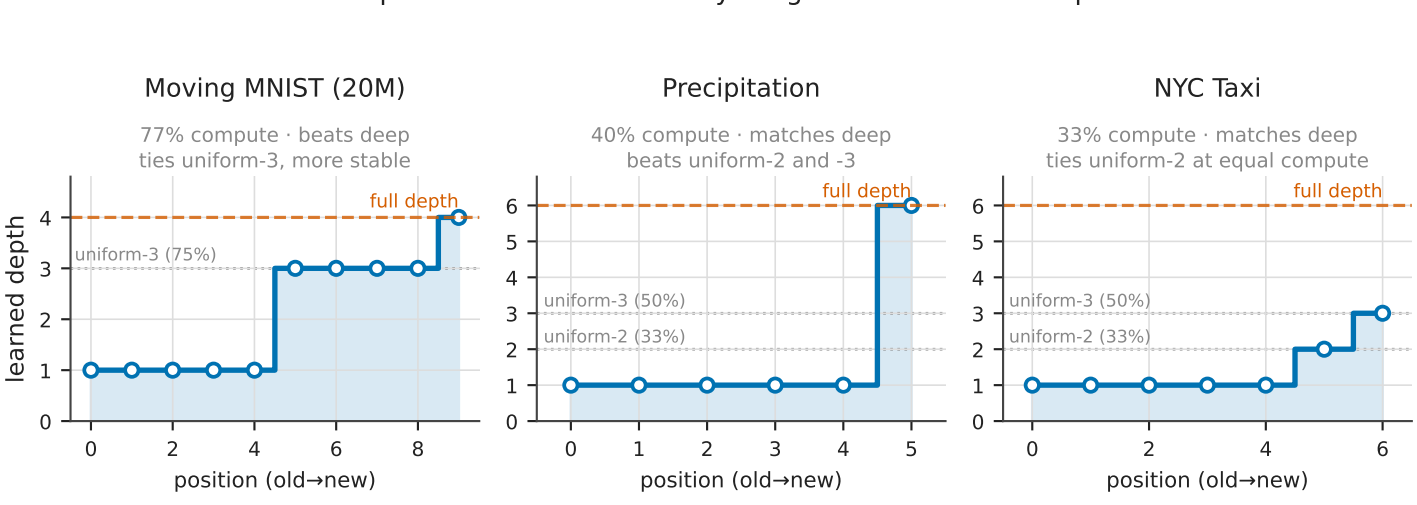
Learned depth schedules are recency-weighted and dataset-dependent
Moving MNIST (20M)
77% compute · beats deep
ties uniform-3, more stable
Precipitation
40% compute · matches deep
beats uniform-2 and -3
NYC Taxi
33% compute · matches deep
ties uniform-2 at equal compute
learned depth
full depth
uniform-3 (75%)
uniform-3 (50%)
uniform-2 (33%)
position (old→new)