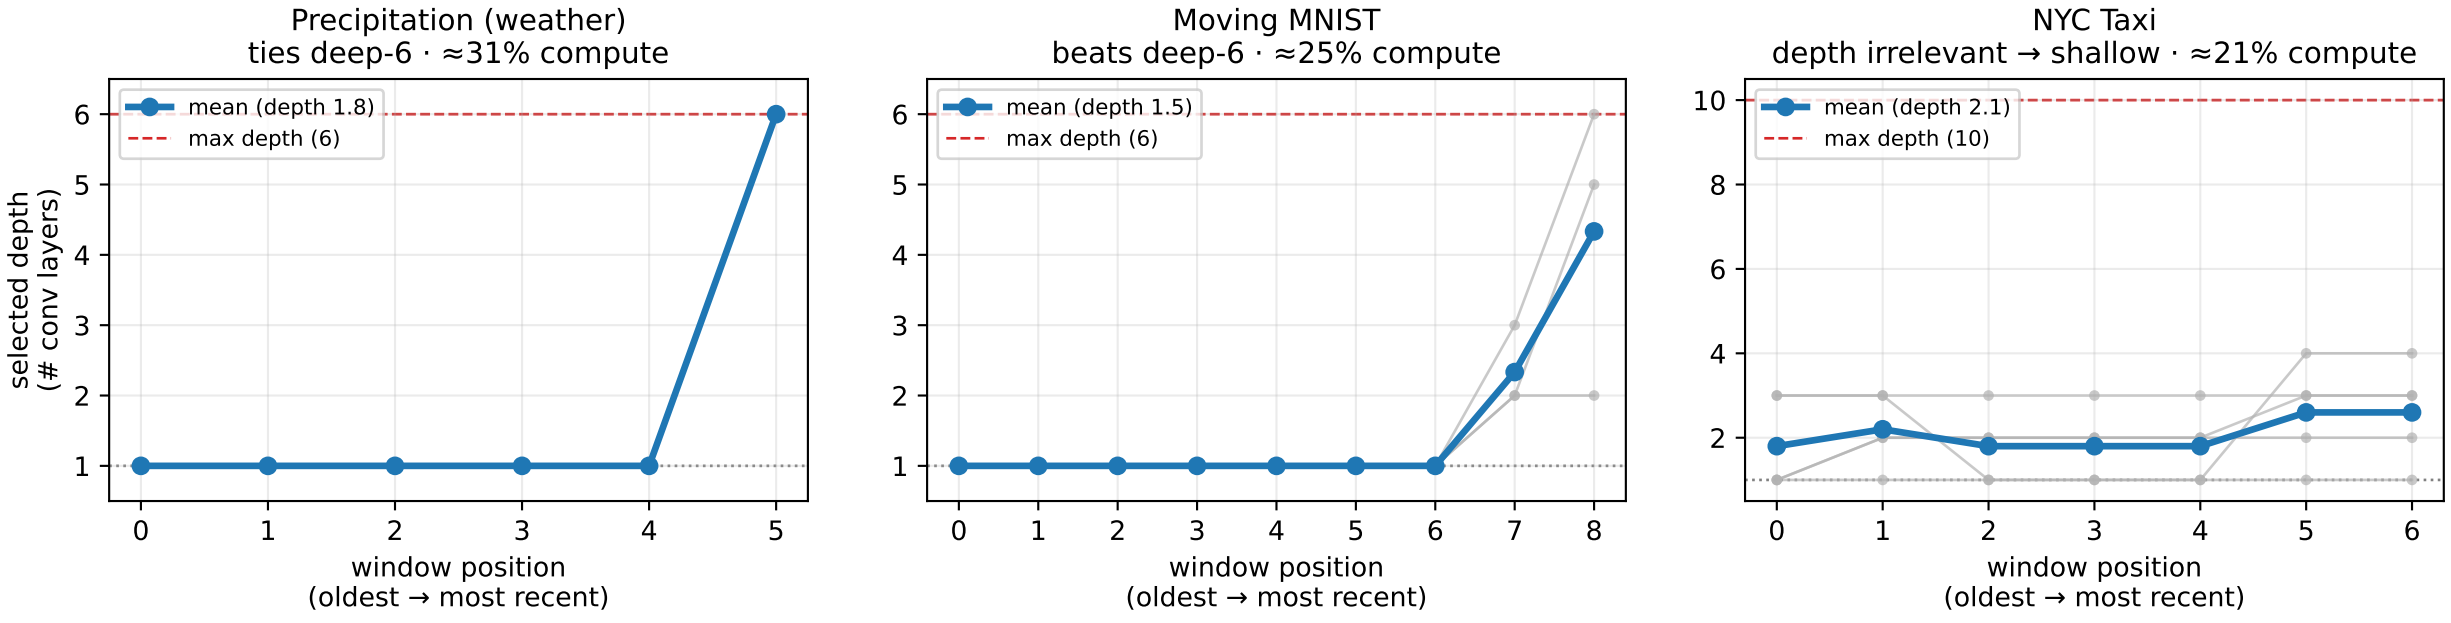
Learned positional depth schedules: the router allocates depth per dataset, no tuning
Precipitation (weather)
ties deep-6 · ≈31% compute
Moving MNIST
beats deep-6 · ≈25% compute
NYC Taxi
depth irrelevant → shallow · ≈21% compute
selected depth
(# conv layers)
mean (depth 1.8)
max depth (6)
mean (depth 1.5)
max depth (6)
mean (depth 2.1)
max depth (10)
window position
(oldest → most recent)
window position
(oldest → most recent)
window position
(oldest → most recent)